

- 🧠 Job Title: Entry-Level AI & Prompt Engineer
- 📍 Location: Remote or Chicago, IL
- 💼 Job Type: Full-time | Entry-Level
- 📁 Department: Artificial Intelligence / NLP

About the Role

We're looking for a motivated and detail-oriented Entry-Level AI & Prompt Engineer to join our growing AI team. This is an opportunity to work hands-on with the latest large language models (LLMs) like GPT-4, Claude, and open-source alternatives. You'll be responsible for crafting, testing, and optimizing prompts that enable LLMs to deliver accurate, safe, and contextually relevant outputs across a variety of use cases.

Whether it's powering conversational AI, enhancing document search with RAG (retrieval-augmented generation), or fine-tuning task-specific behaviors, your contributions will play a key role in shaping intelligent and human-aligned applications.

Key Responsibilities

- Design and refine prompts for LLM-based tasks, including summarization, classification, information extraction, and question answering
- Collaborate with AI engineers, product managers, and UX teams to integrate prompt-based solutions into production systems
- Evaluate and benchmark prompt performance across various models and versions (e.g., OpenAI GPT, Anthropic Claude, Mistral)
- Implement few-shot, zero-shot, and chain-of-thought prompting strategies
- Use tools such as LangChain, LlamaIndex, or vector databases (Pinecone, FAISS) to support hybrid RAG workflows
- Document prompt libraries, workflows, and experiment outcomes
- Stay current on the latest in prompt engineering, LLMs, and AI safety best practices

Qualifications

Required:

- Bachelor's degree (or final semester) in Computer Science, Data Science, Computational Linguistics, or a related field
- Proficiency in Python and familiarity with APIs (especially OpenAI, Anthropic, Cohere, etc.)
- Solid understanding of NLP fundamentals and transformer-based models
- Strong written communication and critical thinking skills
- Ability to analyze model behavior, output quality, and edge cases
- Experience using tools like Jupyter Notebooks, Colab, Git, and version control

Preferred:

- Hands-on experience with LangChain, Hugging Face Transformers, or RAG pipelines
- Prior exposure to prompt tuning, fine-tuning, or LoRA methods
- Open-source contributions, technical blog posts, or public GitHub projects in the LLM space
- Participation in AI hackathons or related competitions (e.g., Zindi, Kaggle)

What We Offer

- Opportunity to work with cutting-edge LLMs in a fast-growing team
- Direct mentorship from AI researchers and engineers
- Access to premium API tokens and prompt testing environments
- Competitive entry-level compensation and benefits
- Flexible work environment with remote options
- Career growth opportunities in NLP and AI R&D

To Apply

Please submit your resume, a brief cover letter, and (if available) links to any relevant GitHub projects, personal blogs, or prompt engineering notebooks.